

# HW5\_key

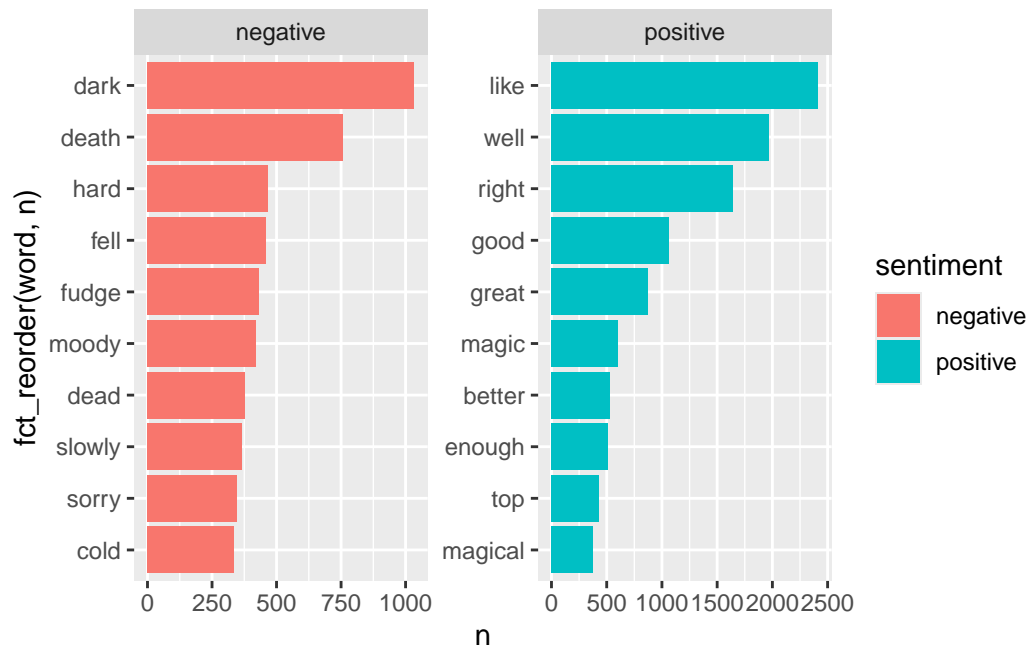
## Harry Potter

The `potter_untidy` dataset includes the text of 7 books of the Harry Potter series by J.K. Rowling. For a brief overview of the books (or movies), see this quote from Wikipedia:

Harry Potter is a series of seven fantasy novels written by British author J. K. Rowling. The novels chronicle the lives of a young wizard, Harry Potter, and his friends Hermione Granger and Ron Weasley, all of whom are students at Hogwarts School of Witchcraft and Wizardry. The main story arc concerns Harry's conflict with Lord Voldemort, a dark wizard who intends to become immortal, overthrow the wizard governing body known as the Ministry of Magic, and subjugate all wizards and Muggles (non-magical people).

1. What words contribute the most to negative and positive sentiment scores? Show a faceted bar plot of the top 10 negative and the top 10 positive words (according to the “bing” lexicon) across the entire series.

```
potter_tidy |>
  inner_join(get_sentiments("bing")) |>
  count(sentiment, word, sort = TRUE) |>
  group_by(sentiment) |>
  top_n(10) |>
  ungroup() |>
  ggplot(aes(x = fct_reorder(word, n), y = n, fill = sentiment)) +
    geom_col() + # makes bar plot where heights = values in the data set
    coord_flip() +
    facet_wrap(~ sentiment, scales = "free")
```



```
# note in warning that enviously appears as both positive and negative
```

2. Find a list of the top 10 words associated with “fear” and with “trust” (according to the “nrc” lexicon) across the entire series.

```
# Check out which words are associated with which sentiment
get_sentiments("nrc") |>
  count(sentiment)
```

```
# A tibble: 10 x 2
  sentiment      n
  <chr>      <int>
1 anger      1245
2 anticipation 837
3 disgust    1056
4 fear       1474
5 joy         687
6 negative    3316
7 positive    2308
8 sadness     1187
9 surprise     532
10 trust      1230
```

```
get_sentiments("nrc") |>
  filter(sentiment == "fear") |>
  inner_join(potter_tidy) |>
  count(word, sort = TRUE)
```

Joining with `by = join\_by(word)`

# A tibble: 887 x 2

	word	n
	<chr>	<int>
1	death	757
2	feeling	391
3	fire	388
4	crouch	297
5	shaking	277
6	scar	276
7	mad	269
8	kill	267
9	elf	259
10	watch	256

# i 877 more rows

```
get_sentiments("nrc") |>
  filter(sentiment == "trust") |>
  inner_join(potter_tidy) |>
  count(word, sort = TRUE)
```

Joining with `by = join\_by(word)`

# A tibble: 676 x 2

	word	n
	<chr>	<int>
1	professor	2006
2	good	1065
3	school	634
4	found	614
5	ministry	576
6	top	434
7	sir	419
8	feeling	391

```
9 lord          391
10 ground       386
# i 666 more rows
```

3. Make a wordcloud for the entire series after removing stop words using the “smart” source.

```
# wordcloud wants a column with words and another column with counts
words <- potter_tidy |>
  anti_join(stop_words) |>
  anti_join(potter_names, join_by(word == firstname)) |>
  anti_join(potter_names, join_by(word == lastname)) |>
  count(word) |>
  arrange(desc(n))

# Note: this will look better in html than in the Plots window in RStudio
wordcloud(
  words = words$word,
  freq = words$n,
  max.words = 100,
  random.order = FALSE,
  rot.per = 0,
  colors = brewer.pal(6, "Dark2")
)
```



```
# See Z's R Tip of the Day for suggestions on options

# Or for even cooler looks, use wordcloud2 in html
# words_df <- words |>
#   slice_head(n = 80) |>
#   data.frame()

# wordcloud2(words_df, size = .35, shape = 'star')
```

4. Create a wordcloud with the top 20 negative words and the top 20 positive words in the Harry Potter series according to the bing lexicon. The words should be sized by their respective counts and colored based on whether their sentiment is positive or negative. (Feel free to be resourceful and creative to color words by a third variable!)

```
pos_neg <- potter_tidy |>
  inner_join(get_sentiments("bing")) |>
  count(sentiment, word, sort = TRUE) |>
  group_by(sentiment) |>
  top_n(20) |>
  ungroup()
```

Joining with `by = join\_by(word)`

Warning in inner\_join(potter\_tidy, get\_sentiments("bing")): Detected an unexpected many-to-many relationship between the variables in the following join:  
 i Row 41432 of `x` matches multiple rows in `y`.  
 i Row 2698 of `y` matches multiple rows in `x`.  
 i If a many-to-many relationship is expected, set `relationship = "many-to-many"` to silence this warning.

Selecting by n


```
wordcloud(
  words = pos_neg$word,
  freq = pos_neg$n,
  random.order = FALSE,
  rot.per = 0,
  ordered.colors = TRUE,
  colors = brewer.pal(6, "Dark2")[factor(pos_neg$sentiment)]
)
```



```
spread(key = sentiment, value = n, fill = 0) |>
mutate(sentiment = positive - negative) |>
ggplot(aes(x = chapter, y = sentiment, fill = title)) +
  geom_col(show.legend = FALSE) +
  facet_wrap(~title, ncol = 2, scales = "free_x")
```



```
potter_tidy |>
mutate(chapter_factor = factor(chapter)) |>
inner_join(get_sentiments("afinn")) |>
group_by(title, chapter_factor) |>
summarise(mean_sentiment = mean(value)) |>
ggplot(aes(x = chapter_factor, y = mean_sentiment,
           fill = title, color = title)) +
  geom_point() +
  geom_hline(yintercept = 0) +
  facet_wrap(~title, ncol = 2)
```



7. Make a faceted bar plot showing the top 10 words that distinguish each book according to the tf-idf statistic.

```
book_word_count <- potter_tidy |>
  count(word, title, sort = TRUE)

book_tfidf <- book_word_count |>
  bind_tf_idf(word, title, n)

book_tfidf |>
  arrange(-tf_idf)
```

# A tibble: 67,845 x 6

	word	title	n	tf	idf	tf_idf
	<chr>	<fct>	<int>	<dbl>	<dbl>	<dbl>
1	slughorn	Half-Blood Prince	335	0.00196	1.25	0.00245
2	umbridge	Order of the Phoenix	496	0.00192	0.847	0.00162
3	bagman	Goblet of Fire	208	0.00108	1.25	0.00136
4	lockhart	Chamber of Secrets	197	0.00231	0.560	0.00129
5	lupin	Prisoner of Azkaban	369	0.00351	0.336	0.00118
6	winky	Goblet of Fire	145	0.000756	1.25	0.000947
7	champions	Goblet of Fire	84	0.000438	1.95	0.000852



```

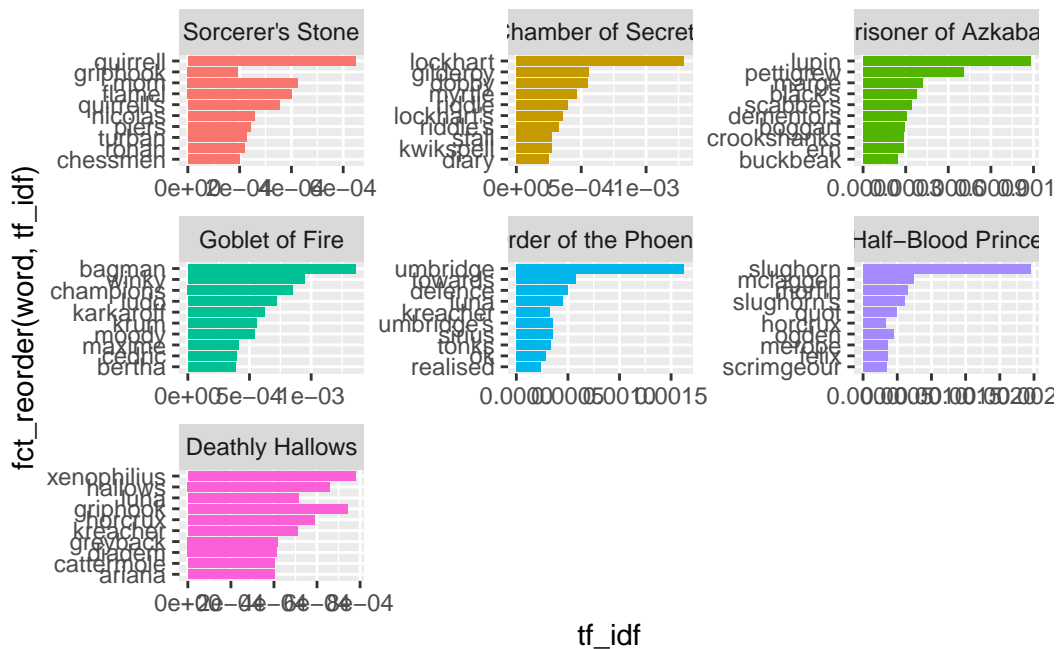
8 xenophilius Deathly Hallows          79 0.000400 1.95 0.000778
9 griphook    Deathly Hallows          117 0.000592 1.25 0.000742
10 mclaggen    Half-Blood Prince        65 0.000379 1.95 0.000738
# i 67,835 more rows

```

```

book_tf_idf |>
  group_by(title) |>
  arrange(desc(tf_idf)) |>
  top_n(10, wt = tf_idf) |>
  ungroup() |>
  ggplot(aes(x = fct_reorder(word, tf_idf), y = tf_idf, fill = title)) +
    geom_col(show.legend = FALSE) +
    coord_flip() +
    facet_wrap(~title, scales = "free")

```



8. Repeat (7) to show the top 10 2-word combinations that distinguish each book.

```

tidy_ngram <- potter_untidy |>
  unnest_tokens(bigram, text, token = "ngrams", n = 2)

bigram_tf_idf <- tidy_ngram |>
  count(title, bigram) |>

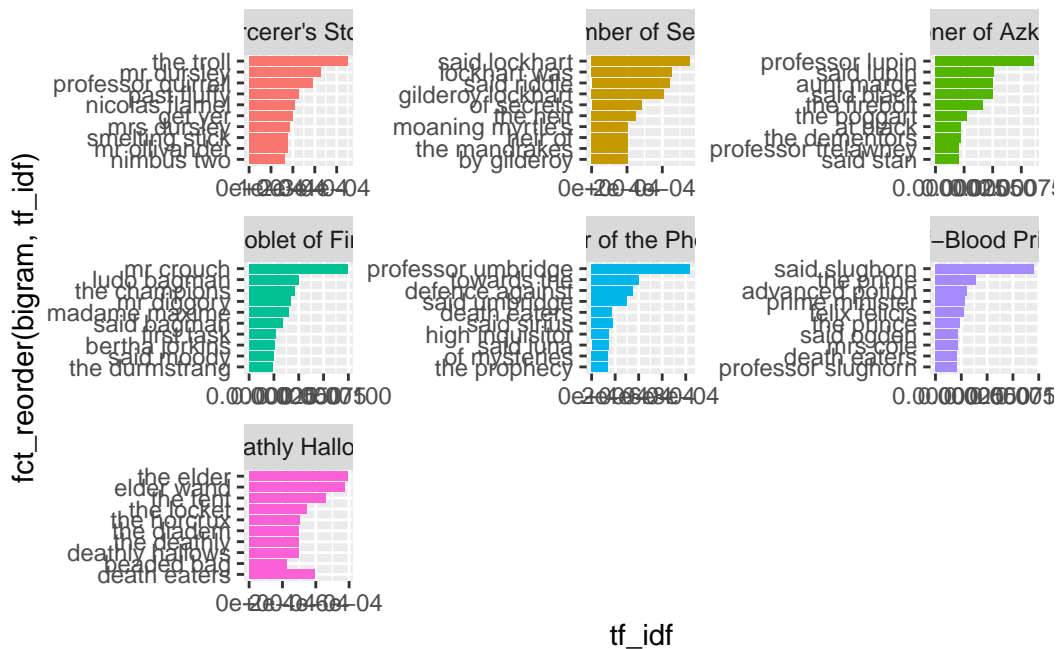
```

```

bind_tf_idf(bigram, title, n) |>
  arrange(desc(tf_idf)) |>
  filter(!is.na(bigram))

bigram_tf_idf |>
  group_by(title) |>
  arrange(desc(tf_idf)) |>
  top_n(10, wt = tf_idf) |>
  ungroup() |>
  ggplot(aes(x = fct_reorder(bigram, tf_idf), y = tf_idf, fill = title)) +
    geom_col(show.legend = FALSE) +
    coord_flip() +
    facet_wrap(~title, scales = "free")

```



- Find which words contributed most in the “wrong” direction using the `afinn` sentiment combined with how often a word appears among all 7 books. Come up with a list of 4 negation words, and for each negation word, illustrate the words associated with the largest “wrong” contributions in a faceted bar plot.

```

afinn <- get_sentiments("afinn")

bigrams_separated <- tidy_ngram |>

```

```

separate(bigram, c("word1", "word2"), sep = " ") |>
count(word1, word2, sort = TRUE) |>
filter(!is.na(word1) & !is.na(word2))

negation_words <- c("not", "no", "never", "without")

negated_words <- bigrams_separated |>
  filter(word1 %in% negation_words) |>
  inner_join(afinn, by = c(word2 = "word")) |>
  arrange(desc(n))

negated_words

```

# A tibble: 379 x 4

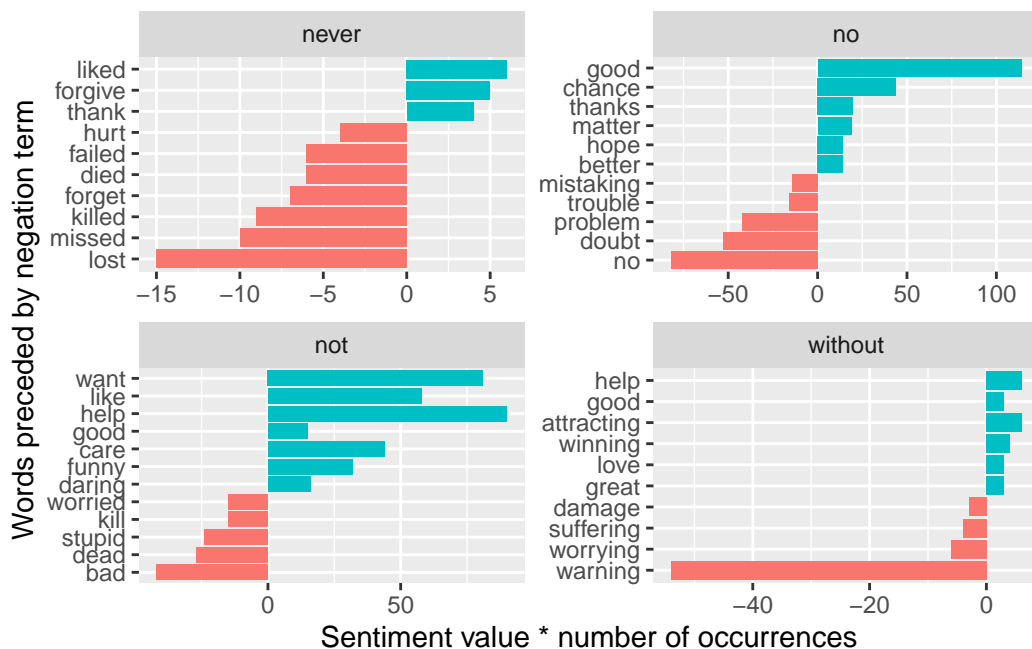
	word1	word2	n	value
	<chr>	<chr>	<int>	<dbl>
1	no	no	82	-1
2	not	want	81	1
3	no	doubt	53	-1
4	not	help	45	2
5	no	good	38	3
6	not	like	29	2
7	no	chance	22	2
8	not	care	22	2
9	no	problem	21	-2
10	no	matter	19	1

# i 369 more rows

```

negated_words |>
  mutate(contribution = n * value) |>
  arrange(desc(abs(contribution))) |>
  group_by(word1) |>
  slice_max(abs(contribution), n = 10) |>
  ungroup() |>
  mutate(word2 = reorder(word2, contribution)) |>
  ggplot(aes(n * value, word2, fill = n * value > 0)) +
    geom_col(show.legend = FALSE) +
    facet_wrap(~ word1, scales = "free") +
    labs(x = "Sentiment value * number of occurrences",
         y = "Words preceded by negation term")

```



10. Select a set of 4 “interesting” terms and then use the Phi coefficient to find and plot the 6 words most correlated with each of your “interesting” words. Start by dividing `potter_tidy` into 80-word sections and then remove names and spells and stop words.

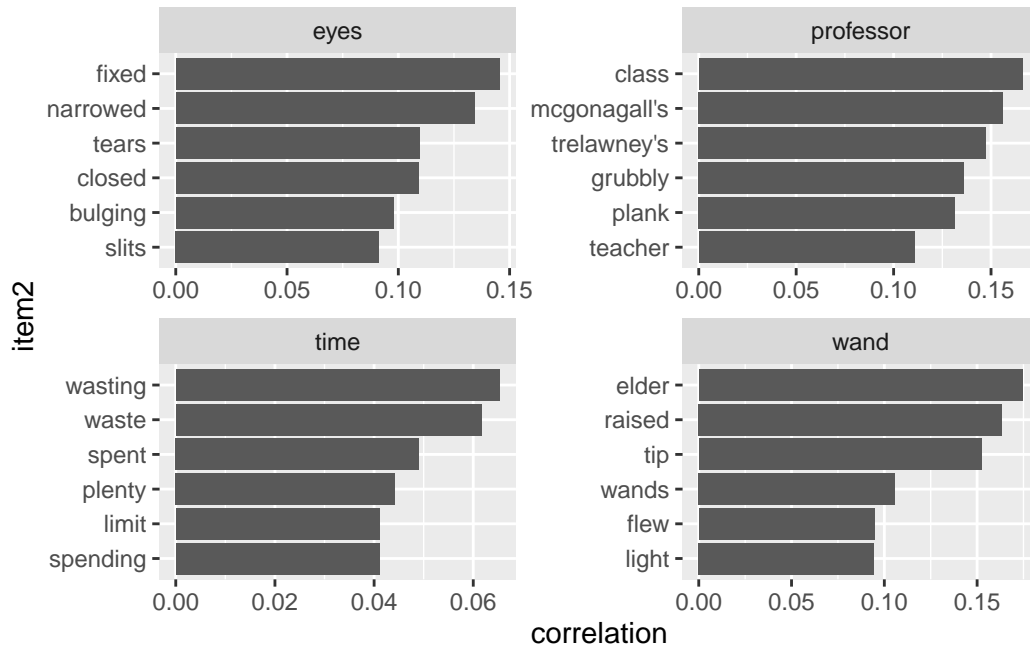
```
library(widyr)

potter_section_words <- potter_tidy |>
  mutate(section = 1 + row_number() %/% 80) |>
  anti_join(potter_names, join_by(word == firstname)) |>
  anti_join(potter_names, join_by(word == lastname)) |>
  anti_join(potter_spells, join_by(word == first_word)) |>
  anti_join(potter_spells, join_by(word == second_word)) |>
  filter(!word %in% stop_words$word,
         !is.na(word))

word_cor <- potter_section_words |>
  group_by(word) |>
  filter(n() >= 10) |>
  pairwise_cor(word, section, sort = TRUE)

# Plot words most associated with a set of interesting words:
word_cor |>
  filter(item1 %in% c("eyes", "professor", "wand", "time")) |>
```

```
group_by(item1) |>
slice_max(correlation, n = 6) |>
ungroup() |>
mutate(item2 = reorder(item2, correlation)) |>
ggplot(aes(item2, correlation)) +
  geom_bar(stat = "identity") +
  facet_wrap(~ item1, scales = "free") +
  coord_flip()
```



11. Create a network graph to visualize the correlations and clusters of words that were found by the `widyr` package in (10).

```
library(igraph)
```

Attaching package: 'igraph'

The following objects are masked from 'package:lubridate':

%--%, union

The following objects are masked from 'package:dplyr':

```
as_data_frame, groups, union
```

The following objects are masked from 'package:purrr':

```
compose, simplify
```

The following object is masked from 'package:tidyr':

```
crossing
```

The following object is masked from 'package:tibble':

```
as_data_frame
```

The following objects are masked from 'package:stats':

```
decompose, spectrum
```

The following object is masked from 'package:base':

```
union
```

```
library(ggraph)
set.seed(1989)

word_cors |>
  filter(correlation > .5) |>
  graph_from_data_frame() |>
  ggraph(layout = "fr") +
    geom_edge_link(aes(edge_alpha = correlation), show.legend = FALSE) +
    geom_node_point(color = "lightblue", size = 5) +
    geom_node_text(aes(label = name), repel = TRUE) +
    theme_void()
```



12. Use LDA to fit a 2-topic model to all 7 Harry Potter books. Be sure to remove names, spells, and stop words before running your topic models. (a) Make a plot to illustrate words with greatest difference between two topics, using log ratio. (b) Print a table with the gamma variable for each document and topic. Based on (a) and (b), can you interpret what the two topics represent?

```
# cast the collection of 3 works as a document-term matrix
library(tm)
```

Loading required package: NLP

```
Attaching package: 'NLP'
```

The following object is masked from 'package:ggplot2':

annotate

```
book_word_count <- potter_tidy |>
  group_by(title, word) |>
  count() |>
```

```

    arrange(desc(n))

seven_books_dtm <- book_word_count |>
  anti_join(potter_names, join_by(word == firstname)) |>
  anti_join(potter_names, join_by(word == lastname)) |>
  anti_join(potter_spells, join_by(word == first_word)) |>
  anti_join(potter_spells, join_by(word == second_word)) |>
  filter(!word %in% stop_words$word,
         !is.na(word)) |>
  cast_dtm(title, word, n)

# set a seed so that the output of the model is predictable
library(topicmodels)
seven_books_lda <- LDA(seven_books_dtm, k = 2, control = list(seed = 1234))
seven_books_lda

```

A LDA\_VEM topic model with 2 topics.

```

seven_books_topics <- tidy(seven_books_lda, matrix = "beta")
seven_books_topics

```

```

# A tibble: 46,830 x 3
  topic term      beta
  <int> <chr>    <dbl>
1     1 1 professor 0.00395
2     2 2 professor 0.00746
3     1 1 wand     0.00352
4     2 2 wand     0.00583
5     1 1 looked  0.00578
6     2 2 looked  0.00766
7     1 1 voice    0.00395
8     2 2 voice    0.00436
9     1 1 time     0.00353
10    2 2 time     0.00623
# i 46,820 more rows

```

```

# Find the most common words within each topic
seven_books_top_terms <- seven_books_topics |>
  group_by(topic) |>
  slice_max(beta, n = 10) |>

```

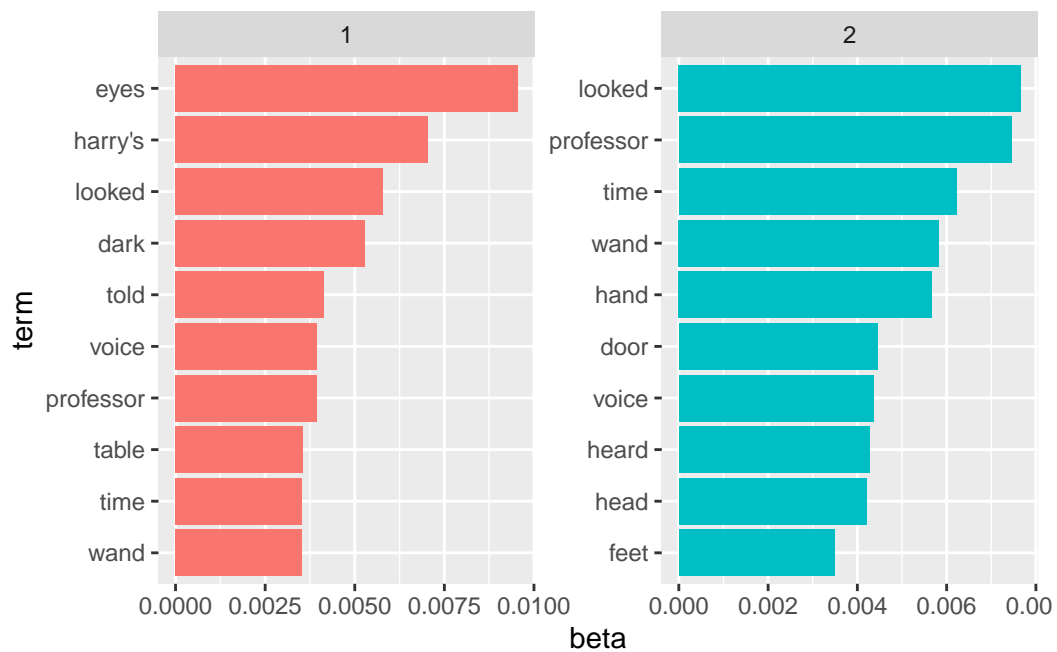


```

ungroup() |>
arrange(topic, -beta)

seven_books_top_terms |>
mutate(term = reorder_within(term, beta, topic)) |>
ggplot(aes(beta, term, fill = factor(topic))) +
  geom_col(show.legend = FALSE) +
  facet_wrap(~ topic, scales = "free") +
  scale_y_reordered()

```



```

# Find words with greatest difference between two topics, using log ratio
beta_wide <- seven_books_topics |>
mutate(topic = paste0("topic", topic)) |>
pivot_wider(names_from = topic, values_from = beta) |>
filter(topic1 > .001 | topic2 > .001) |>
mutate(log_ratio = log2(topic2 / topic1))

beta_wide

```

```

# A tibble: 201 x 4
  term      topic1  topic2 log_ratio
<chr>    <dbl>    <dbl>    <dbl>

```

1	professor	0.00395	0.00746	0.916
2	wand	0.00352	0.00583	0.729
3	looked	0.00578	0.00766	0.407
4	voice	0.00395	0.00436	0.142
5	time	0.00353	0.00623	0.817
6	door	0.00300	0.00445	0.568
7	head	0.00342	0.00420	0.296
8	harry's	0.00705	0.00120	-2.56
9	eyes	0.00954	0.000179	-5.74
10	death	0.00251	0.00189	-0.405

# i 191 more rows

```
beta_wide |>
  arrange(desc(abs(log_ratio))) |>
  slice_max(abs(log_ratio), n = 20) |>
  mutate(term = reorder(term, log_ratio)) |>
  ggplot(aes(log_ratio, term, fill = log_ratio > 0)) +
    geom_col(show.legend = FALSE) +
    labs(x = "Log ratio of Beta values",
         y = "Words in seven works")
```



```
# find the gamma variable for each document and topic
seven_books_documents <- tidy(seven_books_lda, matrix = "gamma")
seven_books_documents
```

```
# A tibble: 14 x 3
  document          topic gamma
  <chr>            <int> <dbl>
1 Sorcerer's Stone      1 0.489
2 Chamber of Secrets    1 0.524
3 Prisoner of Azkaban    1 0.436
4 Goblet of Fire        1 0.501
5 Order of the Phoenix   1 0.466
6 Half-Blood Prince     1 0.457
7 Deathly Hallows       1 0.482
8 Sorcerer's Stone      2 0.511
9 Chamber of Secrets    2 0.476
10 Prisoner of Azkaban   2 0.564
11 Goblet of Fire        2 0.499
12 Order of the Phoenix  2 0.534
13 Half-Blood Prince     2 0.543
14 Deathly Hallows       2 0.518
```

Note: this code could be even better by converting repeated sections into functions!